

# Vaccine Hesitance

## 1. Business Understanding

---

### Context

In 2009, the U.S. conducted the National H1N1 Flu Survey to understand vaccination behaviors during the H1N1 pandemic. This effort mirrors the challenges faced in the COVID-19 era. Understanding the factors that influence vaccine uptake can help public health officials design better vaccine access strategies, communication campaigns, and policy decisions.

### Project Goal

**Predict whether a person received the H1N1 flu vaccine** based on their background, health behavior, beliefs, and socio-demographics.

This project aims to:

- Help public health departments **identify hesitant populations**
  - Guide **targeted awareness campaigns**
  - Support **data-driven resource allocation** (e.g., mobile clinics, outreach)
  - Inform the design of **future vaccination strategies**
- 

### ◆ Key Stakeholders

- **Public health officials** managing vaccination programs
  - **Healthcare providers** engaging with patients and communities
  - **Policy makers** determining communication and funding priorities
  - **Data scientists** supporting health systems and NGOs
- 

### ◆ Business Questions

1. Who is most likely to get vaccinated?
2. What demographic or behavioral features are associated with vaccine acceptance or hesitancy?
3. Can we predict unvaccinated individuals **before** a campaign begins?
4. How can public health campaigns be improved based on these insights?

## 2.Data Understanding

```
In [4]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

### Step 1. Loading the datasets provided

```
In [6]: # Training set features
df_train_features = pd.read_csv('data/training_set_features.csv')
df_train_features.head()
```

```
Out[6]:
```

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavioral_av
0	0	1.0	0.0	0.0	
1	1	3.0	2.0	0.0	
2	2	1.0	1.0	0.0	
3	3	1.0	1.0	0.0	
4	4	2.0	1.0	0.0	

5 rows × 36 columns



```
In [7]: # Getting a general feel of what the df_train_features looks like

print(f"The shape of df_train_features is :\n", df_train_features.shape)
print()
print(f"Some of the important features of the dataset", df_train_features.info())
print()
# print(f"Snapshot of what the numeric data looks like \n", feature_labels_df.describe())
# Check for missing values
print("Missing Values in Features:")
print(df_train_features.isnull().sum().sort_values(ascending=False).head(10))
```

The shape of df\_train\_features is :  
(26707, 36)

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 26707 entries, 0 to 26706

Data columns (total 36 columns):

#	Column	Non-Null Count	Dtype
0	respondent_id	26707 non-null	int64
1	h1n1_concern	26615 non-null	float64
2	h1n1_knowledge	26591 non-null	float64
3	behavioral_antiviral_meds	26636 non-null	float64
4	behavioral_avoidance	26499 non-null	float64
5	behavioral_face_mask	26688 non-null	float64
6	behavioral_wash_hands	26665 non-null	float64
7	behavioral_large_gatherings	26620 non-null	float64
8	behavioral_outside_home	26625 non-null	float64
9	behavioral_touch_face	26579 non-null	float64
10	doctor_recc_h1n1	24547 non-null	float64
11	doctor_recc_seasonal	24547 non-null	float64
12	chronic_med_condition	25736 non-null	float64
13	child_under_6_months	25887 non-null	float64
14	health_worker	25903 non-null	float64
15	health_insurance	14433 non-null	float64
16	opinion_h1n1_vacc_effective	26316 non-null	float64
17	opinion_h1n1_risk	26319 non-null	float64
18	opinion_h1n1_sick_from_vacc	26312 non-null	float64
19	opinion_seas_vacc_effective	26245 non-null	float64
20	opinion_seas_risk	26193 non-null	float64
21	opinion_seas_sick_from_vacc	26170 non-null	float64
22	age_group	26707 non-null	object
23	education	25300 non-null	object
24	race	26707 non-null	object
25	sex	26707 non-null	object
26	income_poverty	22284 non-null	object
27	marital_status	25299 non-null	object
28	rent_or_own	24665 non-null	object
29	employment_status	25244 non-null	object
30	hhs_geo_region	26707 non-null	object
31	census_msa	26707 non-null	object
32	household_adults	26458 non-null	float64
33	household_children	26458 non-null	float64
34	employment_industry	13377 non-null	object
35	employment_occupation	13237 non-null	object

dtypes: float64(23), int64(1), object(12)

memory usage: 7.3+ MB

Some of the important features of the dataset None

Missing Values in Features:

employment_occupation	13470
employment_industry	13330
health_insurance	12274
income_poverty	4423
doctor_recc_seasonal	2160
doctor_recc_h1n1	2160
rent_or_own	2042

```

employment_status      1463
marital_status          1408
education               1407
dtype: int64

```

```

In [8]: # Looking into the Labels dataset
df_train_labels = pd.read_csv('data/training_set_labels.csv')
df_train_labels.head()

```

```

Out[8]:
   respondent_id  h1n1_vaccine  seasonal_vaccine
0              0              0                  0
1              1              0                  1
2              2              0                  0
3              3              0                  1
4              4              0                  0

```

```

In [9]: # Getting a general of what the df_train_features Looks Like

print(f"The shape of df_train_labels is :\n", df_train_labels.shape)
print()
print(f"Some of the important features of the dataset", df_train_labels.info())
print()
# print(f"Snapshot of what the numeric data looks like \n", df_train_labels.describe())
# Check for missing values
print("Missing Values in Features:")
print(df_train_labels.isnull().sum().sort_values(ascending=False).head(10))

```

The shape of df\_train\_labels is :

```
(26707, 3)
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26707 entries, 0 to 26706
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   respondent_id    26707 non-null  int64
1   h1n1_vaccine     26707 non-null  int64
2   seasonal_vaccine 26707 non-null  int64
dtypes: int64(3)
memory usage: 626.1 KB
Some of the important features of the dataset None

```

```

Missing Values in Features:
respondent_id      0
h1n1_vaccine       0
seasonal_vaccine   0
dtype: int64

```

```

In [10]: # Test set features
# This particular dataset will act as the used data our models

```

```
df_test_features = pd.read_csv('data/test_set_features.csv')
df_test_features.head()
```

Out[10]:

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavioral_av
--	---------------	--------------	----------------	---------------------------	---------------

0	26707	2.0	2.0	0.0	
---	-------	-----	-----	-----	--

1	26708	1.0	1.0	0.0	
---	-------	-----	-----	-----	--

2	26709	2.0	2.0	0.0	
---	-------	-----	-----	-----	--

3	26710	1.0	1.0	0.0	
---	-------	-----	-----	-----	--

4	26711	3.0	1.0	1.0	
---	-------	-----	-----	-----	--

5 rows × 36 columns



```
In [11]: # Getting a general of what the df_test_features Looks Like

print(f"The shape of df_train_labels is :\n", df_test_features.shape)
print()
print(f"Some of the important features of the dataset", df_test_features.info())
print()
# print(f"Snapshot of what the numeric data looks like \n", df_test_features.describe())
# Check for missing values
print("Missing Values in Features:")
print(df_test_features.isnull().sum().sort_values(ascending=False).head(10))
```

The shape of df\_train\_labels is :  
(26708, 36)

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 26708 entries, 0 to 26707

Data columns (total 36 columns):

#	Column	Non-Null Count	Dtype
0	respondent_id	26708 non-null	int64
1	h1n1_concern	26623 non-null	float64
2	h1n1_knowledge	26586 non-null	float64
3	behavioral_antiviral_meds	26629 non-null	float64
4	behavioral_avoidance	26495 non-null	float64
5	behavioral_face_mask	26689 non-null	float64
6	behavioral_wash_hands	26668 non-null	float64
7	behavioral_large_gatherings	26636 non-null	float64
8	behavioral_outside_home	26626 non-null	float64
9	behavioral_touch_face	26580 non-null	float64
10	doctor_recc_h1n1	24548 non-null	float64
11	doctor_recc_seasonal	24548 non-null	float64
12	chronic_med_condition	25776 non-null	float64
13	child_under_6_months	25895 non-null	float64
14	health_worker	25919 non-null	float64
15	health_insurance	14480 non-null	float64
16	opinion_h1n1_vacc_effective	26310 non-null	float64
17	opinion_h1n1_risk	26328 non-null	float64
18	opinion_h1n1_sick_from_vacc	26333 non-null	float64
19	opinion_seas_vacc_effective	26256 non-null	float64
20	opinion_seas_risk	26209 non-null	float64
21	opinion_seas_sick_from_vacc	26187 non-null	float64
22	age_group	26708 non-null	object
23	education	25301 non-null	object
24	race	26708 non-null	object
25	sex	26708 non-null	object
26	income_poverty	22211 non-null	object
27	marital_status	25266 non-null	object
28	rent_or_own	24672 non-null	object
29	employment_status	25237 non-null	object
30	hhs_geo_region	26708 non-null	object
31	census_msa	26708 non-null	object
32	household_adults	26483 non-null	float64
33	household_children	26483 non-null	float64
34	employment_industry	13433 non-null	object
35	employment_occupation	13282 non-null	object

dtypes: float64(23), int64(1), object(12)

memory usage: 7.3+ MB

Some of the important features of the dataset None

Missing Values in Features:

employment_occupation	13426
employment_industry	13275
health_insurance	12228
income_poverty	4497
doctor_recc_seasonal	2160
doctor_recc_h1n1	2160
rent_or_own	2036

```
employment_status      1471
marital_status          1442
education               1407
dtype: int64
```

## Step 2: Data Cleaning and Preprocessing

Before proceeding to analysis, the raw data was thoroughly reviewed and cleaned to ensure consistency, reliability, and readiness for modeling.

### Objectives

- Prepare the dataset for bivariate and multivariate analysis.
- Standardize formats and handle missing or inconsistent entries.
- Implement reusable, object-oriented code for cleaning workflows.

### Cleaning Workflow

The following steps were taken:

#### 1. Initial Exploration

- Loaded and previewed the dataset structure.
- Checked shape, column names, and data types.

#### 2. Handling Missing Values and Inconsistencies

- Identified and quantified missing values across columns.
- Filled missing values:
  - Numeric columns → median
  - Categorical columns → mode
  - Text/object columns → "Unknown"

#### 3. Type Conversion

- Converted appropriate `float` and `object` columns with limited unique values to `category` dtype to optimize memory and modeling.

#### 4. Code Reusability

- Implemented a custom class `VaccineDataCleaner` using Object-Oriented Programming (OOP) principles to:
  - Merge features and labels
  - Clean and prepare the dataset
  - Return a clean, analysis-ready DataFrame

---

The cleaned dataset ( `df_clean` ) is now ready for:

- Exploratory Data Analysis (EDA)

- Bivariate and multivariate visualizations
- Feature selection and model training

In [13]: `from data_cleaning import VaccineDataCleaner`

In [14]: `# implementing VaccineDataCleaner on the dataset`

```
cleaner = VaccineDataCleaner(df_train_features, df_train_labels)
cleaner.fill_missing()
cleaner.convert_types()

# returning the final cleaned dataframe
df_clean = cleaner.get_cleaned_data()
df_clean.head()
```

Out[14]:

	<b>respondent_id</b>	<b>h1n1_concern</b>	<b>h1n1_knowledge</b>	<b>behavioral_antiviral_meds</b>	<b>behavioral_av</b>
--	----------------------	---------------------	-----------------------	----------------------------------	----------------------

<b>0</b>	0	1.0	0.0	0.0	
----------	---	-----	-----	-----	--

<b>1</b>	1	3.0	2.0	0.0	
----------	---	-----	-----	-----	--

<b>2</b>	2	1.0	1.0	0.0	
----------	---	-----	-----	-----	--

<b>3</b>	3	1.0	1.0	0.0	
----------	---	-----	-----	-----	--

<b>4</b>	4	2.0	1.0	0.0	
----------	---	-----	-----	-----	--

5 rows × 38 columns



In [15]: `# Getting a general of what the df_clean looks like`

```
print(f"The shape of df_train_labels is :\n", df_clean.shape)
print()
print(f"Some of the important features of the dataset", df_clean.info())
print()
# print(f"Snapshot of what the numeric data looks like \n", df_clean.describe())
# Check for missing values
print("Missing Values in Features:")
print(df_clean.isnull().sum().sort_values(ascending=False).head(10))
```



The shape of df\_train\_labels is :  
(26707, 38)

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 26707 entries, 0 to 26706

Data columns (total 38 columns):

#	Column	Non-Null Count	Dtype
0	respondent_id	26707 non-null	int64
1	h1n1_concern	26707 non-null	category
2	h1n1_knowledge	26707 non-null	category
3	behavioral_antiviral_meds	26707 non-null	category
4	behavioral_avoidance	26707 non-null	category
5	behavioral_face_mask	26707 non-null	category
6	behavioral_wash_hands	26707 non-null	category
7	behavioral_large_gatherings	26707 non-null	category
8	behavioral_outside_home	26707 non-null	category
9	behavioral_touch_face	26707 non-null	category
10	doctor_recc_h1n1	26707 non-null	category
11	doctor_recc_seasonal	26707 non-null	category
12	chronic_med_condition	26707 non-null	category
13	child_under_6_months	26707 non-null	category
14	health_worker	26707 non-null	category
15	health_insurance	26707 non-null	category
16	opinion_h1n1_vacc_effective	26707 non-null	category
17	opinion_h1n1_risk	26707 non-null	category
18	opinion_h1n1_sick_from_vacc	26707 non-null	category
19	opinion_seas_vacc_effective	26707 non-null	category
20	opinion_seas_risk	26707 non-null	category
21	opinion_seas_sick_from_vacc	26707 non-null	category
22	age_group	26707 non-null	category
23	education	26707 non-null	category
24	race	26707 non-null	category
25	sex	26707 non-null	category
26	income_poverty	26707 non-null	category
27	marital_status	26707 non-null	category
28	rent_or_own	26707 non-null	category
29	employment_status	26707 non-null	category
30	hhs_geo_region	26707 non-null	category
31	census_msa	26707 non-null	category
32	household_adults	26707 non-null	category
33	household_children	26707 non-null	category
34	employment_industry	26707 non-null	object
35	employment_occupation	26707 non-null	object
36	h1n1_vaccine	26707 non-null	int64
37	seasonal_vaccine	26707 non-null	int64

dtypes: category(33), int64(3), object(2)

memory usage: 1.9+ MB

Some of the important features of the dataset None

Missing Values in Features:

respondent_id	0
h1n1_concern	0
h1n1_knowledge	0
behavioral_antiviral_meds	0
behavioral_avoidance	0

```
behavioral_face_mask      0
behavioral_wash_hands     0
behavioral_large_gatherings 0
behavioral_outside_home   0
behavioral_touch_face     0
dtype: int64
```

## Step 3: Exploratory Data Analysis (EDA)


This section begins the EDA process using a custom `VaccineEDA` class to explore variable distributions, relationships, and patterns relevant to vaccine uptake.


```
In [17]: from EDA import VaccineEDA
```

We can start by looking at the distribution of the two target variables, from the `df_train_labels`

```
In [19]: eda = VaccineEDA(df_clean)

eda.plot_bar('h1n1_vaccine'), eda.plot_bar('seasonal_vaccine')
```





Out[19]: (None, None)

From the two plots above we can see that almost half of the sampled population recieved the seaosol flu vaccine but less than 30% received the H1N1 vaccine. The seasonal vaccine has class balance while the H1N1 is quite imbalanced.

- *Next we can assess whether the two target variables are independent*

```
In [21]: pd.crosstab(
          df_clean["h1n1_vaccine"], df_clean["seasonal_vaccine"], margins=True, normalize
        )
```

```
Out[21]: seasonal_vaccine    0    1    All
          h1n1_vaccine
          -----
          0  0.497810  0.289737  0.787546
          1  0.036582  0.175871  0.212454
          All 0.534392  0.465608  1.000000
```

## Ordinal Encoding and Correlation Analysis

To prepare the data for correlation analysis and modeling:

- **Applied Ordinal Encoding**

All remaining categorical columns were converted to numeric format using `OrdinalEncoder`, which assigns integer values to each category. This is important because correlation requires numeric input.

- **Dropped object columns**

These columns are non-numeric and not suitable for correlation analysis.

```
In [23]: from sklearn.preprocessing import OrdinalEncoder

encoder = OrdinalEncoder()
df_encoded = df_clean.copy()
categorical_cols = df_encoded.select_dtypes(include='category').columns
df_encoded[categorical_cols] = encoder.fit_transform(
    df_encoded[categorical_cols])

# intention is to drop the columns employment_industry and employment_occupation
# since they may prove difficult to encode
df_encoded = df_encoded.drop(columns=['employment_industry', 'employment_occupation'])
```

### Select Top 10 Correlated Features

The top 10 features with the highest absolute correlation to the target features selected for further exploration.

```
In [25]: # Now compute correlation with the target = h1n1_vaccine
correlations = df_encoded.corr()
target_col = 'h1n1_vaccine'
correlations[target_col].drop(labels=[target_col]).sort_values(ascending=False).head()
```

```
Out[25]: doctor_recc_h1n1          0.394086
seasonal_vaccine          0.377143
opinion_h1n1_risk         0.320580
opinion_h1n1_vacc_effective 0.267352
opinion_seas_risk         0.255874
doctor_recc_seasonal      0.218976
opinion_seas_vacc_effective 0.177799
health_worker             0.168056
h1n1_concern              0.121574
h1n1_knowledge            0.117771
Name: h1n1_vaccine, dtype: float64
```

- ***The most significant factor for getting the H1N1 vaccine is doctor's recommendation followed closely by whether the individual has already gotten the seasonal flu vaccine***

```
In [27]: # Now compute correlation with the target = sesonal_vaccine
correlations = df_encoded.corr()
target_col1 = 'seasonal_vaccine'
correlations[target_col1].drop(labels=[target_col1]).sort_values(ascending=False).head()
```

```
Out[27]: opinion_seas_risk          0.386916
h1n1_vaccine          0.377143
doctor_recc_seasonal  0.360696
opinion_seas_vacc_effective  0.358869
age_group             0.277454
opinion_h1n1_risk      0.215650
opinion_h1n1_vacc_effective  0.203187
doctor_recc_h1n1       0.198560
chronic_med_condition  0.169465
h1n1_concern           0.154488
Name: seasonal_vaccine, dtype: float64
```

### Next Step – Crosstab Analysis

For each of the top correlated features:

- A crosstab will be created showing the relationship between feature values and vaccine uptake ( **0** = not vaccinated, **1** = vaccinated).
- These visualizations will help identify which features may be strong predictors for modeling.


```
In [29]: # Step 1: Create dictionary of top correlated features per target
top_corr_dict = {}
```

```
for target in ['h1n1_vaccine', 'seasonal_vaccine']:
    top_corr = df_encoded.corr()[target] \
        .drop(labels=[target]) \
        .abs() \
        .sort_values(ascending=False) \
        .head(10) \
        .index.tolist()


    top_corr_dict[target] = top_corr
```

```
# Step 2: Loop through and plot using
for target, feature_list in top_corr_dict.items():
    print(f" Crosstab plots for target: `{target}`")
    for feature in feature_list:
        eda.crosstab_plot(feature, target)
```


```
Crosstab plots for target: `h1n1_vaccine`
h1n1_vaccine          0          1
doctor_recc_h1n1
0.0                   0.868773  0.131227
1.0                   0.467641  0.532359
```



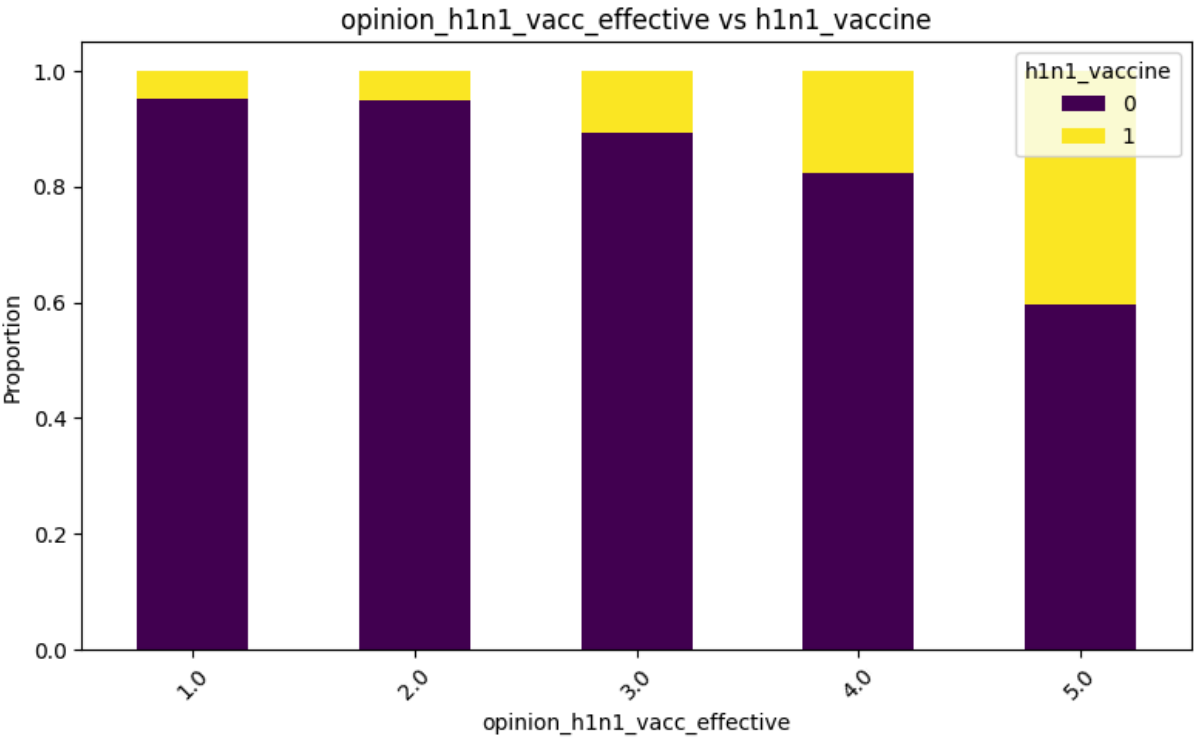
h1n1_vaccine	0	1
seasonal_vaccine		
0	0.931544	0.068456
1	0.622276	0.377724




h1n1_vaccine	0	1
opinion_h1n1_risk		
1.0	0.911660	0.088340
2.0	0.830018	0.169982
3.0	0.826321	0.173679
4.0	0.607898	0.392102
5.0	0.489143	0.510857




h1n1_vaccine	0	1
opinion_h1n1_vacc_effective		
1.0	0.952596	0.047404
2.0	0.948870	0.051130
3.0	0.894347	0.105653
4.0	0.823008	0.176992
5.0	0.595172	0.404828



h1n1_vaccine	0	1
opinion_seas_risk		
1.0	0.915300	0.084700
2.0	0.839987	0.160013
3.0	0.799114	0.200886
4.0	0.693971	0.306029
5.0	0.600406	0.399594

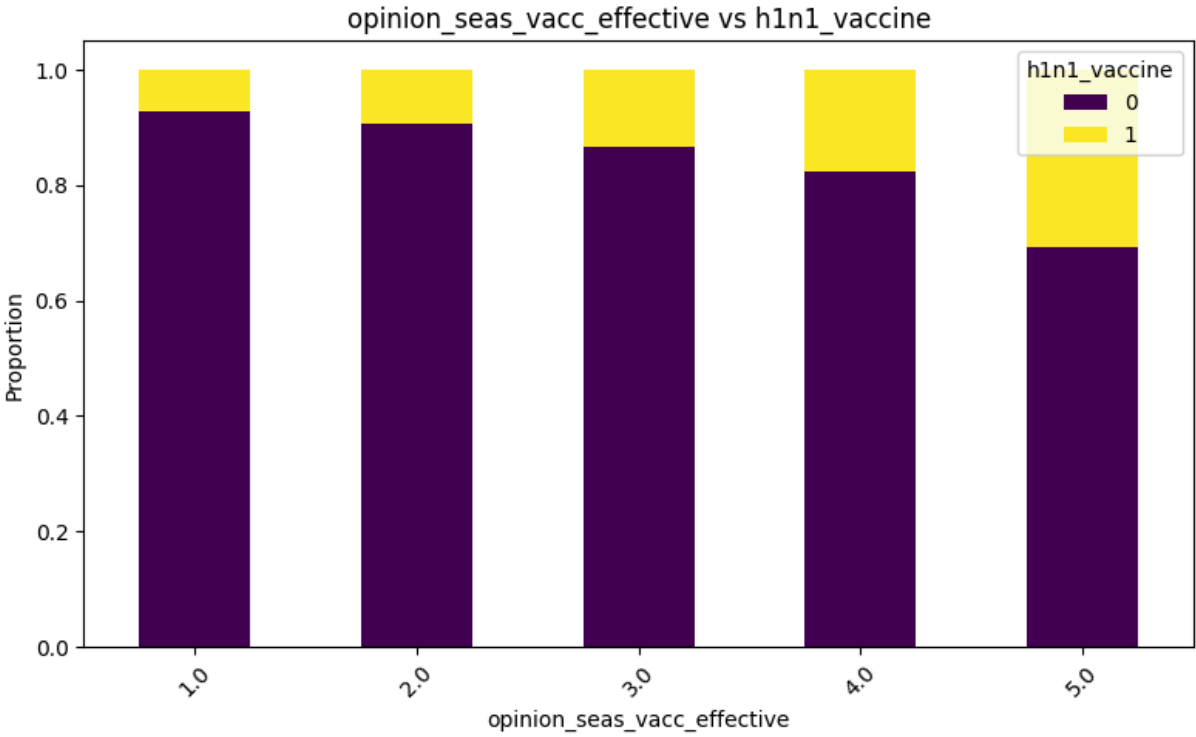


h1n1_vaccine	0	1
doctor_recc_seasonal		
0.0	0.846613	0.153387
1.0	0.651717	0.348283

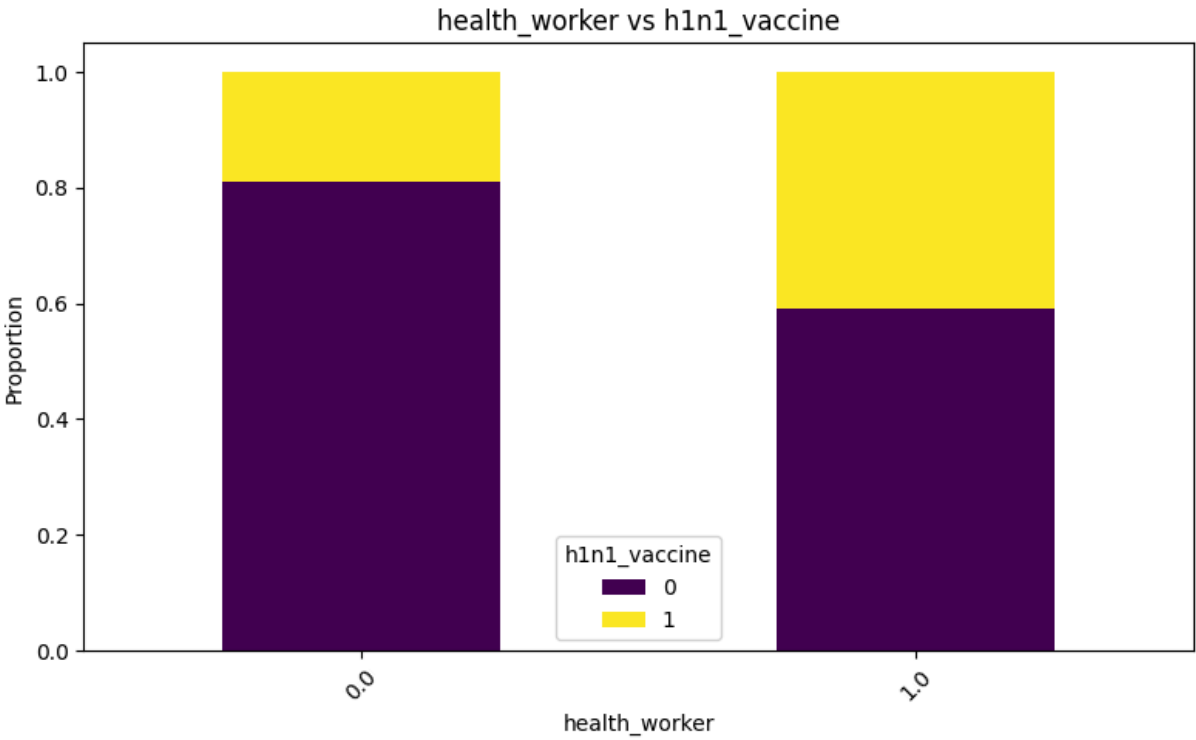




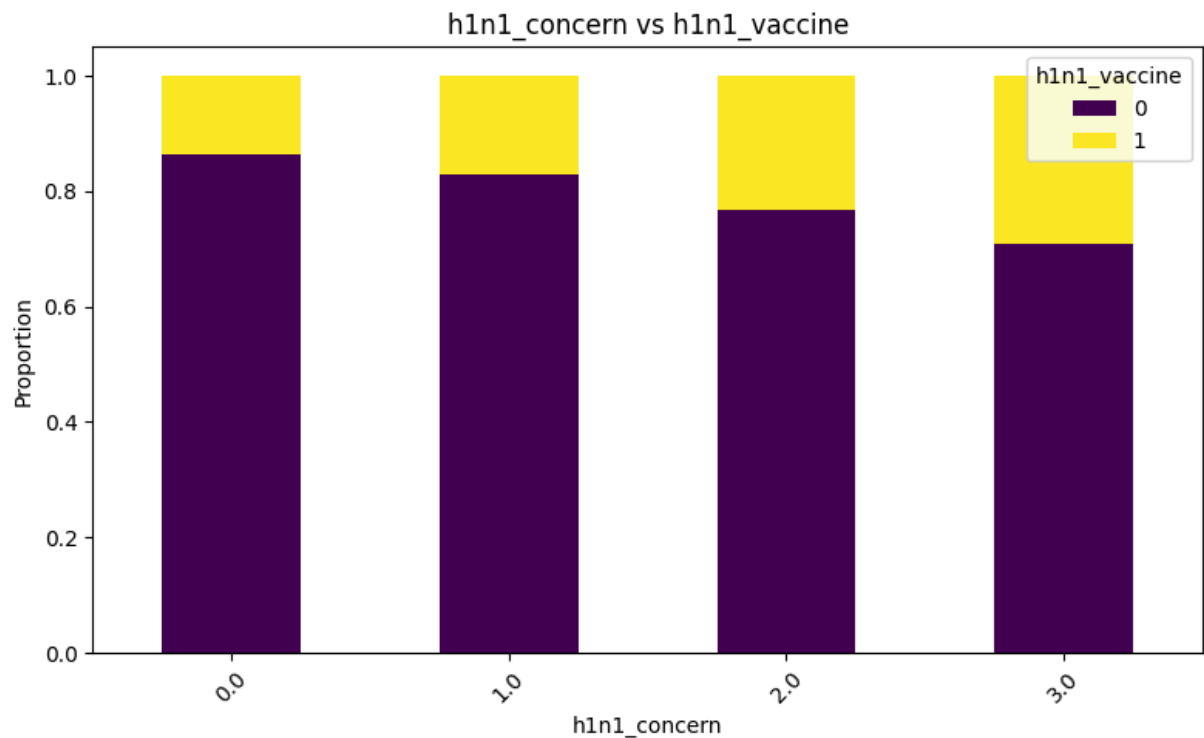
h1n1_vaccine	0	1
opinion_seas_vacc_effective		
1.0	0.927109	0.072891
2.0	0.907072	0.092928
3.0	0.866776	0.133224
4.0	0.822843	0.177157
5.0	0.691567	0.308433



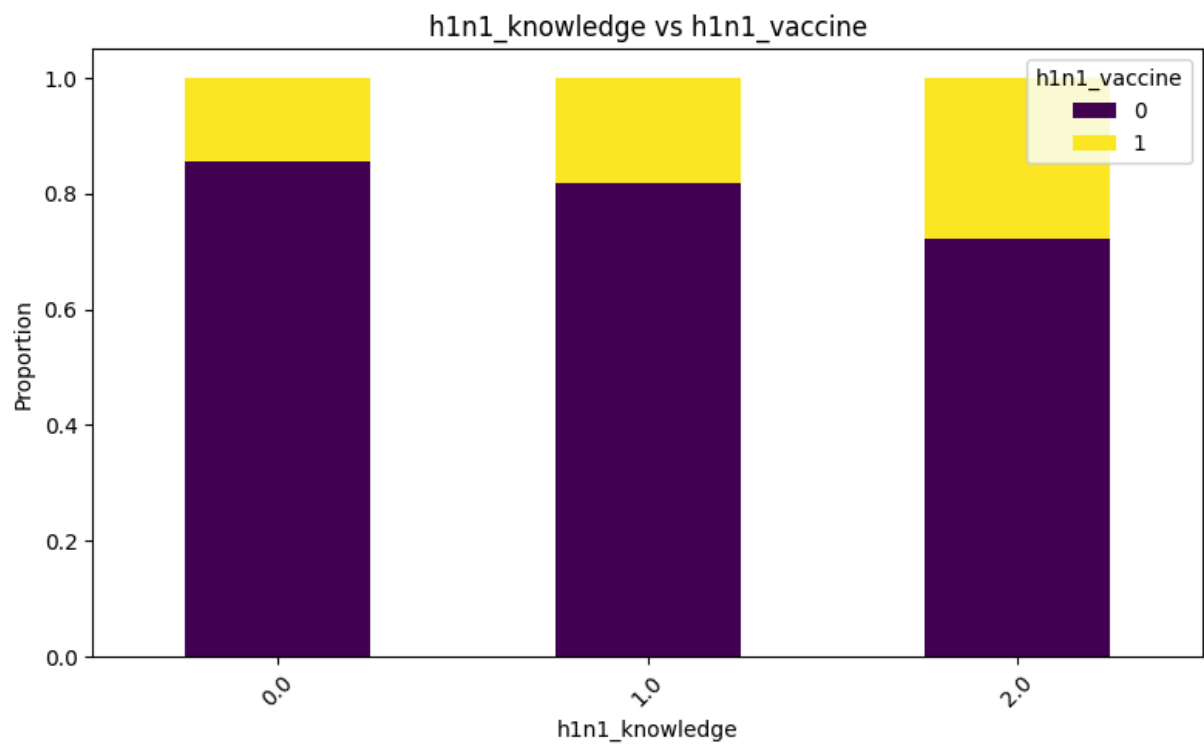
h1n1_vaccine	0	1
health_worker		
0.0	0.811534	0.188466
1.0	0.590548	0.409452



h1n1_vaccine	0	1
h1n1_concern		
0.0	0.864381	0.135619
1.0	0.828652	0.171348
2.0	0.766664	0.233336
3.0	0.707907	0.292093




h1n1_vaccine	0	1
h1n1_knowledge		
0.0	0.855946	0.144054
1.0	0.818200	0.181800
2.0	0.721935	0.278065

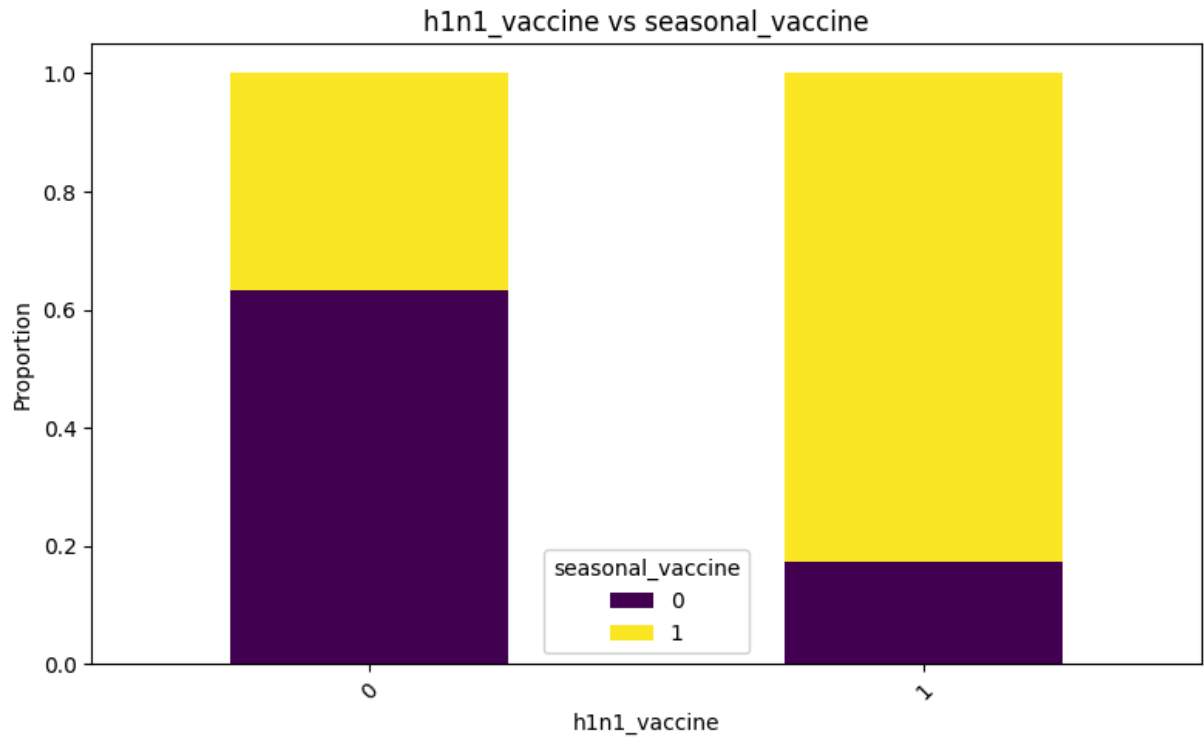


Crosstab plots for target: `seasonal\_vaccine`


seasonal_vaccine	0	1
opinion_seas_risk		
1.0	0.790593	0.209407
2.0	0.625898	0.374102
3.0	0.443131	0.556869
4.0	0.336566	0.663434
5.0	0.255240	0.744760



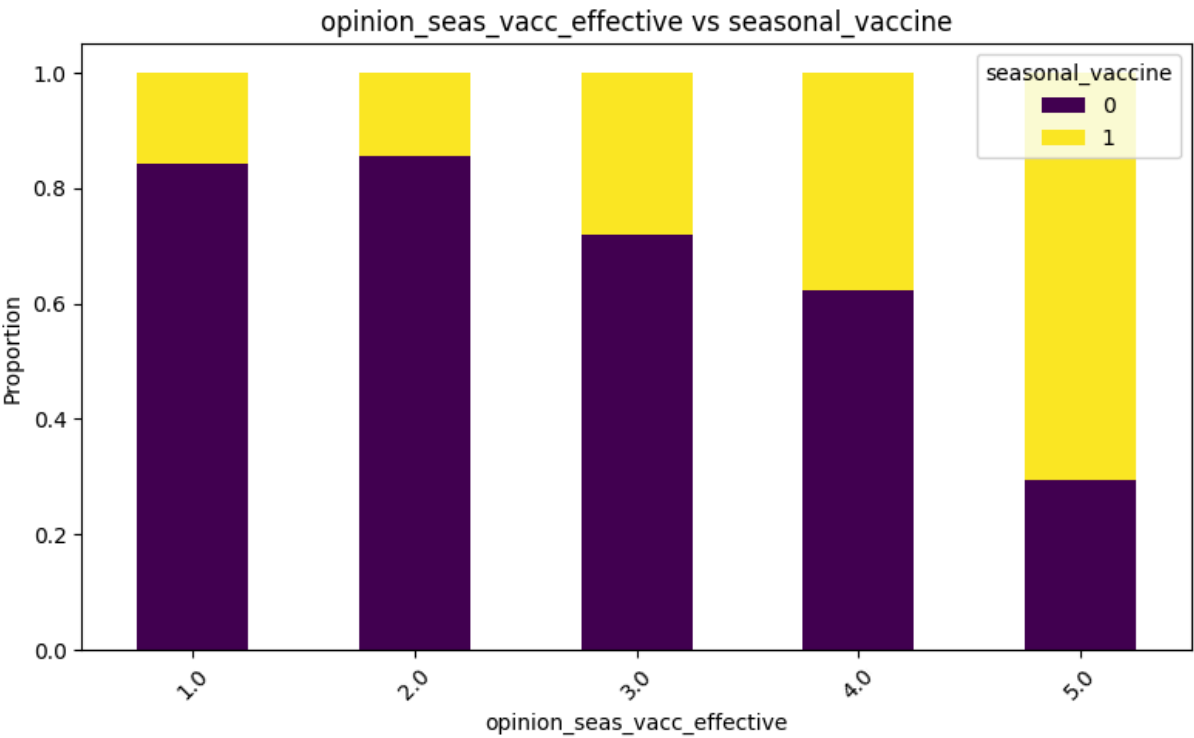
seasonal_vaccine	0	1
h1n1_vaccine		
0	0.632102	0.367898
1	0.172189	0.827811




seasonal_vaccine	0	1
doctor_recc_seasonal		
0.0	0.653038	0.346962
1.0	0.261552	0.738448



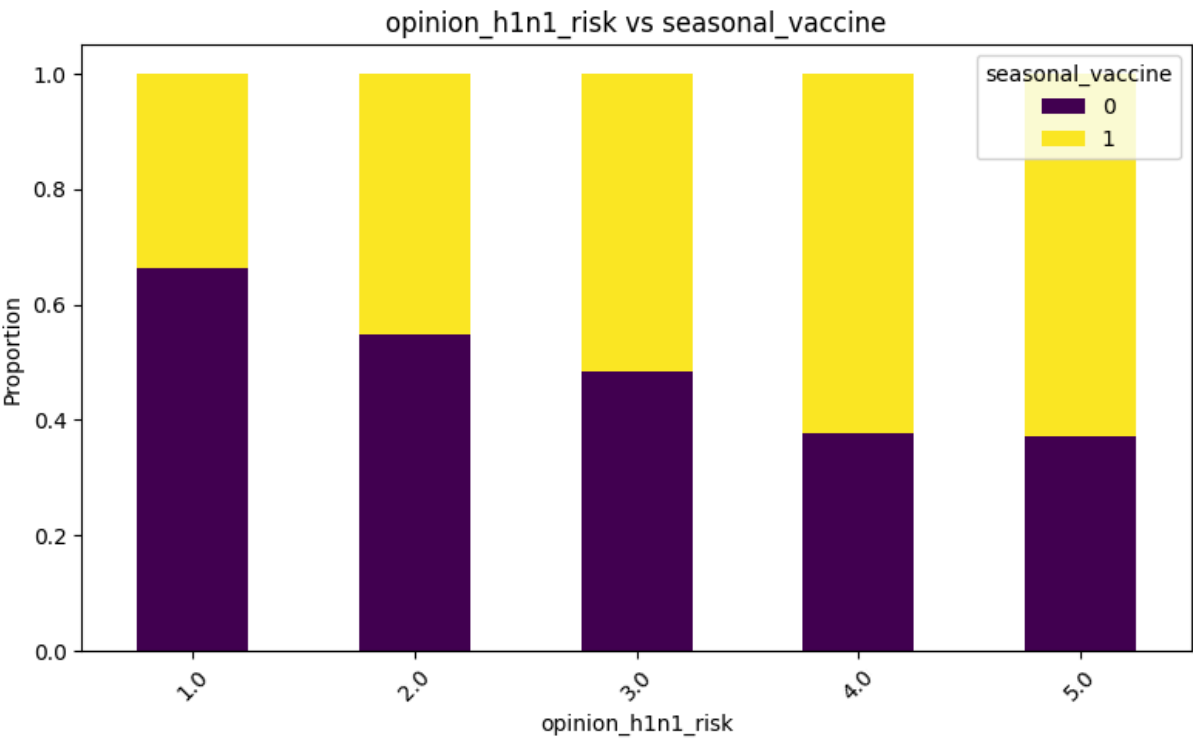
seasonal_vaccine	0	1
opinion_seas_vacc_effective		
1.0	0.841933	0.158067
2.0	0.854488	0.145512
3.0	0.720395	0.279605
4.0	0.623439	0.376561
5.0	0.295297	0.704703




seasonal_vaccine	0	1
age_group		
18 - 34 Years	0.715436	0.284564
35 - 44 Years	0.637474	0.362526
45 - 54 Years	0.598702	0.401298
55 - 64 Years	0.488765	0.511235
65+ Years	0.326319	0.673681



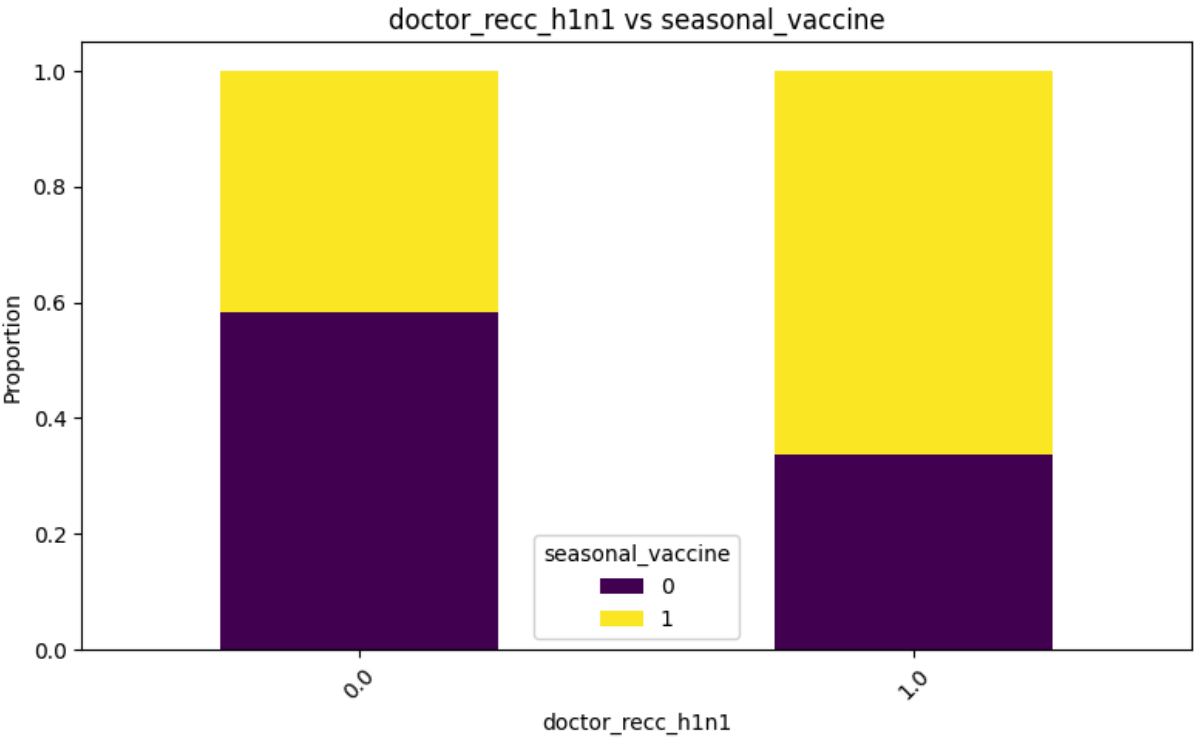
seasonal_vaccine	0	1
opinion_h1n1_risk		
1.0	0.664332	0.335668
2.0	0.547492	0.452508
3.0	0.483438	0.516562
4.0	0.376159	0.623841
5.0	0.373143	0.626857



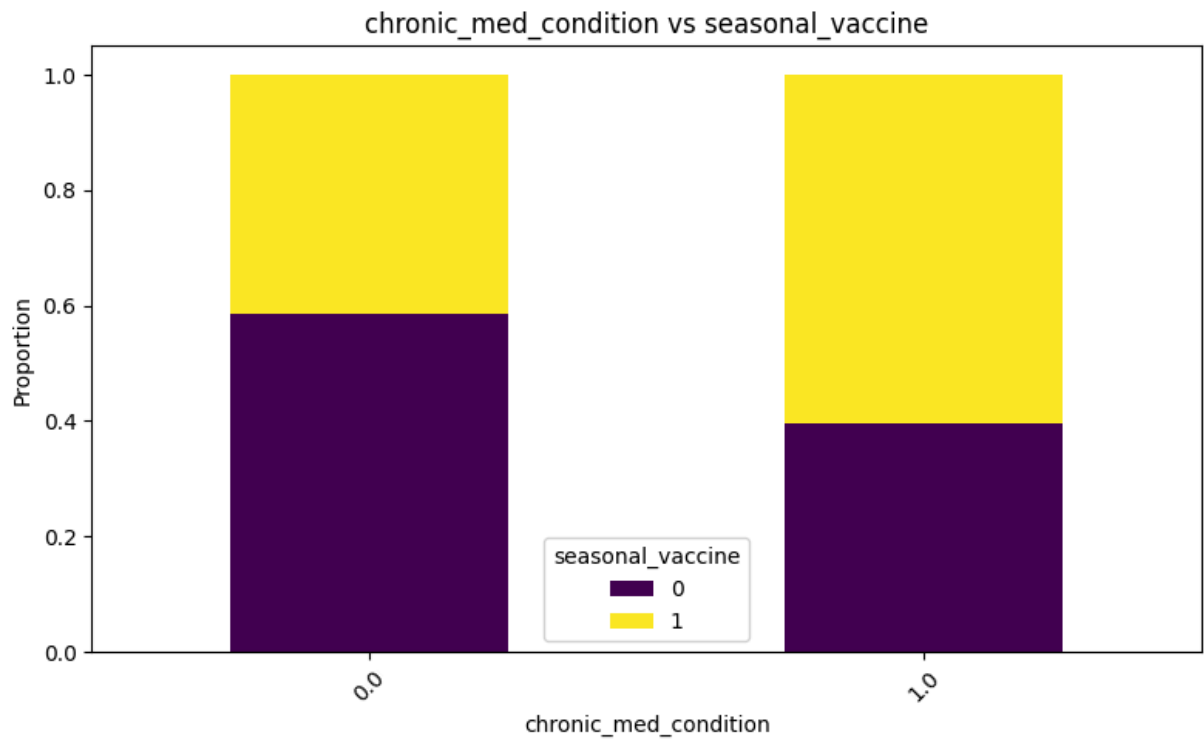
seasonal_vaccine	0	1
opinion_h1n1_vacc_effective		
1.0	0.788939	0.211061
2.0	0.743272	0.256728
3.0	0.577387	0.422613
4.0	0.553835	0.446165
5.0	0.387664	0.612336



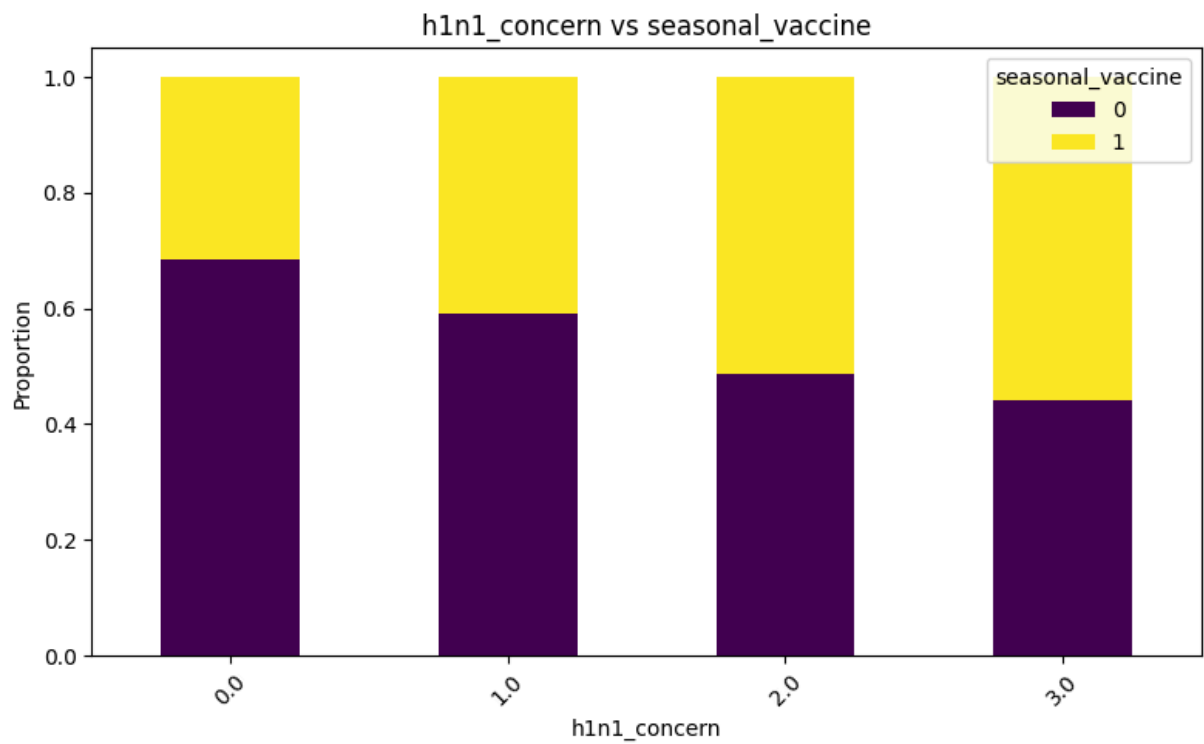
seasonal_vaccine	0	1
doctor_recc_h1n1		
0.0	0.584300	0.415700
1.0	0.337833	0.662167



seasonal_vaccine	0	1
chronic_med_condition		
0.0	0.586187	0.413813
1.0	0.396433	0.603567



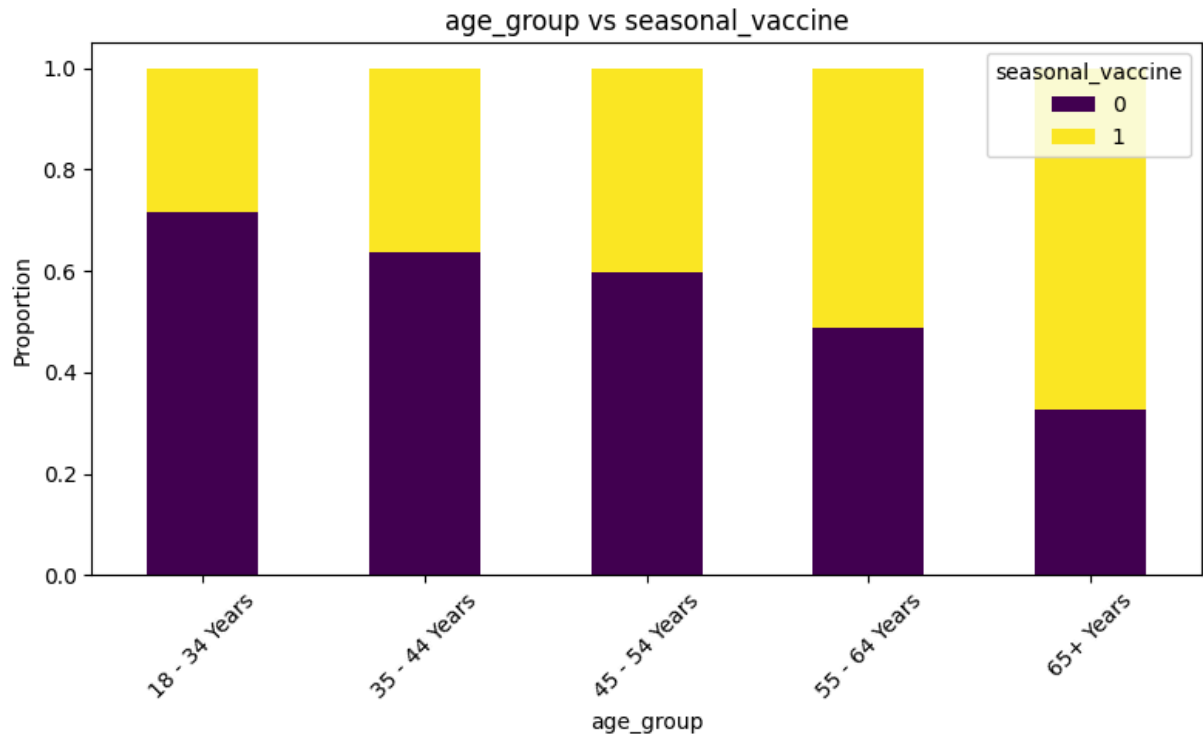
seasonal_vaccine	0	1
h1n1_concern		
0.0	0.683859	0.316141
1.0	0.590090	0.409910
2.0	0.486079	0.513921
3.0	0.440427	0.559573



```
In [30]: eda.crosstab_plot('age_group', 'seasonal_vaccine')
```



seasonal_vaccine	0	1
age_group		
18 - 34 Years	0.715436	0.284564
35 - 44 Years	0.637474	0.362526
45 - 54 Years	0.598702	0.401298
55 - 64 Years	0.488765	0.511235
65+ Years	0.326319	0.673681



## Step 4: Modeling – Predicting Vaccine Uptake

With the data cleaned, encoded, and explored, we now move into the **modeling phase**.

In this section, we aim to:

- Define our features (  $X$  ) and target (  $y$  )
- Split the dataset into training and test sets
- Train baseline classification models (e.g., Logistic Regression, Decision Tree)
- Evaluate model performance using metrics such as **accuracy, precision, recall, and F1-score**
- Use confusion matrices and classification reports for deeper insight

We will begin with simple baseline models and progressively improve using:

- Hyperparameter tuning
- Cross-validation
- Feature engineering

Our goal: Build an effective model to predict whether a person received the **H1N1 vaccine** or the **Seasonal flu vaccine**, based on their characteristics, behaviors, and opinions.

```
In [32]: from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.linear_model import LinearRegression, LogisticRegression, RidgeCV, Lasso
from sklearn.model_selection import cross_val_score
from sklearn.preprocessing import PolynomialFeatures
from sklearn.cluster import KMeans
from sklearn.metrics import confusion_matrix
from sklearn.tree import DecisionTreeClassifier

# imports from VaccineModel class
from modelling import VaccineModel
# from sklearn.linear_model import LogisticRegression
# from sklearn.tree import DecisionTreeClassifier
# from sklearn.pipeline import Pipeline
# from sklearn.preprocessing import StandardScaler
# from sklearn.model_selection import train_test_split
# from sklearn.metrics import (
#     accuracy_score, precision_score, recall_score,
#     f1_score, confusion_matrix, classification_report
# )
```

## Baseline Modeling: Logistic Regression vs. Decision Tree

To establish a benchmark, we start training two baseline models on the encoded dataset:

- **Logistic Regression:** A linear model suitable for binary classification
- **Decision Tree:** A non-linear, rule-based model that can capture interactions

Here we use baseline parameters for both and evaluate performance using accuracy, precision, recall, and F1-score. This helps us compare model complexity vs. predictive performance in future iterations.

## Baseline Modelling for H1N1

```
In [35]: # Logistic regression is set as default in VaccineModel class

h1n1_lr_base = VaccineModel(df_encoded, target='h1n1_vaccine', scale=True)
h1n1_lr_base.split_data()
h1n1_lr_base.train_model()
h1n1_lr_base.evaluate()
```

Evaluation for `h1n1\_vaccine` using Pipeline:

Accuracy: 0.8356420816173717  
 Precision: 0.6848920863309352  
 Recall: 0.4193832599118943  
 F1 Score: 0.5202185792349727

Confusion Matrix:

```
[[3988  219]
 [ 659  476]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.86	0.95	0.90	4207
1	0.68	0.42	0.52	1135
accuracy			0.84	5342
macro avg	0.77	0.68	0.71	5342
weighted avg	0.82	0.84	0.82	5342

```
In [36]: # Decision Tree
tree_model = DecisionTreeClassifier(max_depth=5, random_state=42)

h1n1_tree_base = VaccineModel(df_encoded, target='h1n1_vaccine', model=tree_model)
h1n1_tree_base.split_data()
h1n1_tree_base.train_model()
h1n1_tree_base.evaluate()
```

Evaluation for `h1n1\_vaccine` using DecisionTreeClassifier:

Accuracy: 0.8264694870834893  
 Precision: 0.6268292682926829  
 Recall: 0.452863436123348  
 F1 Score: 0.5258312020460358

Confusion Matrix:

```
[[3901  306]
 [ 621  514]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.86	0.93	0.89	4207
1	0.63	0.45	0.53	1135
accuracy			0.83	5342
macro avg	0.74	0.69	0.71	5342
weighted avg	0.81	0.83	0.82	5342

### Confusion Matrix Visualiation

```
In [38]: def plot_confusion_matrices(y_true_lr, y_pred_lr, y_true_tree, y_pred_tree, labels=
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
```

```

# Logistic Regression
cm_lr = confusion_matrix(y_true_lr, y_pred_lr, labels=labels)
sns.heatmap(cm_lr, annot=True, fmt='d', cmap='Blues', ax=axes[0])
axes[0].set_title("Logistic Regression")
axes[0].set_xlabel("Predicted")
axes[0].set_ylabel("Actual")

# Decision Tree
cm_tree = confusion_matrix(y_true_tree, y_pred_tree, labels=labels)
sns.heatmap(cm_tree, annot=True, fmt='d', cmap='Greens', ax=axes[1])
axes[1].set_title("Decision Tree")
axes[1].set_xlabel("Predicted")
axes[1].set_ylabel("Actual")


plt.tight_layout()
plt.show()

```

```

In [39]: # plot confusion matrix for H1N1 vaccine compaing LogReg and decisiontree
plot_confusion_matrices(
    y_true_lr=h1n1_lr_base.y_test,
    y_pred_lr=h1n1_lr_base.y_pred,
    y_true_tree=h1n1_tree_base.y_test,
    y_pred_tree=h1n1_tree_base.y_pred
)

```



### Baseline model Evaluation for H1N1 vaccine

Metric	Logistic Regression	Decision Tree
Accuracy	83.6%	82.6%
Precision	0.685	0.627
Recall	0.419	0.453
F1 Score	0.520	0.526

### Classification Report for Logistic Regression:

- Class 0 (Not Vaccinated): Precision 0.86 | Recall 0.95 | F1 Score 0.90
- Class 1 (Vaccinated): Precision 0.68 | Recall 0.42 | F1 Score 0.52

**Confusion Matrix** | | Predicted 0 | Predicted 1 | |-----|-----|-----| |  
**Actual 0** | 3988 | 306 | | **Actual 1** | 621 | 514 |

**Classification Report for Decision Tree:**

- Class 0 (Not Vaccinated): Precision 0.86 | Recall 0.93 | F1 Score 0.89
- Class 1 (Vaccinated): Precision 0.63 | Recall 0.45 | F1 Score 0.53

**Confusion Matrix**

	Predicted 0	Predicted 1
Actual 0	3901	219
Actual 1	659	476

**Key Insights**

- Logistic Regression has slightly **higher precision and accuracy**, but **lower recall**.
- Decision Tree offers a **better balance** between precision and recall, resulting in a slightly better F1 score.
- Both models struggle with predicting the vaccinated group (class 1), suggesting the need for addressing class imbalance or testing other algorithms

**Baseline Modelling for Seasonal Vaccine**

```
In [46]: # Logistic regression is set as default in VaccineModel class

seasonal_lr_base = VaccineModel(df_encoded, target='seasonal_vaccine', scale=True)
seasonal_lr_base.split_data()
seasonal_lr_base.train_model()
seasonal_lr_base.evaluate()
```

Evaluation for `seasonal\_vaccine` using Pipeline:

Accuracy: 0.7789217521527517  
 Precision: 0.7831743278404163  
 Recall: 0.7261761158021713  
 F1 Score: 0.7535989985395368

Confusion Matrix:

```
[[2355  500]
 [ 681 1806]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.78	0.82	0.80	2855
1	0.78	0.73	0.75	2487
accuracy			0.78	5342
macro avg	0.78	0.78	0.78	5342
weighted avg	0.78	0.78	0.78	5342

```
In [48]: # Decision Tree
tree_model1 = DecisionTreeClassifier(max_depth=5, random_state=42)

seasonal_tree_base = VaccineModel(df_encoded, target='seasonal_vaccine', model=tree_model1)
seasonal_tree_base.split_data()
seasonal_tree_base.train_model()
seasonal_tree_base.evaluate()
```

Evaluation for `seasonal\_vaccine` using DecisionTreeClassifier:

Accuracy: 0.7523399475851741  
 Precision: 0.7886904761904762  
 Recall: 0.6393244873341375  
 F1 Score: 0.7061958694203864

Confusion Matrix:

```
[[2429  426]
 [ 897 1590]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.73	0.85	0.79	2855
1	0.79	0.64	0.71	2487
accuracy			0.75	5342
macro avg	0.76	0.75	0.75	5342
weighted avg	0.76	0.75	0.75	5342


```
In [50]: # plot confusion matrix for seasnal vaccine compaing LogReg and decisiontree

plot_confusion_matrices(
    y_true_lr=seasonal_lr_base.y_test,
```

```

y_pred_lr=seasonal_lr_base.y_pred,
y_true_tree=seasonal_tree_base.y_test,
y_pred_tree=seasonal_tree_base.y_pred
)

```



## Key Insights and Reommendations

This analysis of the 2009 H1N1 Flu Survey provided key insights into factors that influence vaccine uptake for both H1N1 and seasonal flu. The findings are valuable for guiding future public health strategies and improving vaccine campaign outcomes.

### Key Insights

- **Doctor Recommendations Matter**

Respondents who received a recommendation from a healthcare provider were **significantly more likely** to get vaccinated. This highlights the crucial role of **trusted messengers** in influencing public behavior.

- **Perceptions of Risk and Effectiveness Drive Behavior**

Higher **perceived risk** of flu and greater **confidence in vaccine effectiveness** were positively correlated with vaccine uptake. Conversely, fears of **side effects** were associated with hesitancy.

- **Behavioral Factors Provide Predictive Signals**

Preventive actions such as **mask-wearing, handwashing, and avoiding crowds** were often aligned with vaccine acceptance. These behavioral cues can be early indicators of openness to vaccination.

- **Sociodemographic Patterns are Uneven**

Vaccine uptake varied by **age group, education level, and employment sector**. Tailoring campaigns to different demographics can improve reach and effectiveness.

- **Model Performance is Promising**

Both Logistic Regression and Decision Tree classifiers achieved over **82% accuracy**. While Logistic Regression performed slightly better in terms of recall and precision for H1N1 predictions, Decision Trees were easier to interpret and showed comparable F1 scores.

---

## Recommendations for Stakeholders

- **Public Health Officials**

Target individuals with low vaccine uptake using behavioral and demographic signals. Strengthen community-based strategies and deploy trained communicators.

- **Healthcare Providers**

Actively discuss vaccine benefits and risks with patients. Provider recommendation is a key driver of acceptance.

- **Policy Makers**

Invest in educational campaigns that correct misconceptions. Prioritize **health literacy** initiatives and fund **data-driven outreach programs**.

- **Data Scientists & NGOs**

Build predictive tools into health information systems to **flag high-risk groups** before campaigns launch. Use interpretable models to guide real-time decisions.

---

By leveraging these insights, vaccination programs can become more strategic, equitable, and impactful.